

Factorization Machines (FM)- novel but proved and promising approach

Evgeniy Marinov

July 2020

Introduction to FM

- ▶ In 2010, Steffen Rendle (currently a senior research scientist at Google), introduced a seminal paper [Rendle - FM].
- ▶ Many years have passed since such an impactful algorithm has been introduced in the world of ML /Recommender Systems/.
- ▶ FMs are a new model class which combines the advantages of Support Vector Machines(SVM)/polynomial regression and factorization models.

Kaggle Competitions

Kaggle Competitions won with Factorization Machines:

- ▶ Criteo's CTR on display ads contest - 2014

https:

[//www.kaggle.com/c/criteo-display-ad-challenge](https://www.kaggle.com/c/criteo-display-ad-challenge)

<http://www.csie.ntu.edu.tw/~r01922136/>

[kaggle-2014-criteo.pdf](#)

Kaggle Competitions

Kaggle Competitions won with Factorization Machines:

- ▶ Criteo's CTR on display ads contest - 2014

`https:`

`//www.kaggle.com/c/criteo-display-ad-challenge`

`http://www.csie.ntu.edu.tw/~r01922136/`

`kaggle-2014-criteo.pdf`

- ▶ Avazu's CTR prediction context

`https://www.kaggle.com/c/avazu-ctr-prediction`

`http://www.csie.ntu.edu.tw/~r01922136/slides/`

`kaggle-avazu.pdf`

Kaggle Competitions

Kaggle Competitions won with Factorization Machines:

- ▶ Criteo's CTR on display ads contest - 2014

https:

[//www.kaggle.com/c/criteo-display-ad-challenge](https://www.kaggle.com/c/criteo-display-ad-challenge)

<http://www.csie.ntu.edu.tw/~r01922136/>

[kaggle-2014-criteo.pdf](#)

- ▶ Avazu's CTR prediction context

<https://www.kaggle.com/c/avazu-ctr-prediction>

<http://www.csie.ntu.edu.tw/~r01922136/slides/>

[kaggle-avazu.pdf](#)

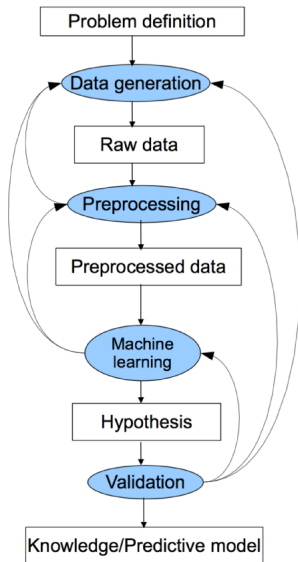
In both competitions the LogLoss function for evaluation of the classifiers has been used - the aim is to **minimize** the Logarithmic Loss.

- ▶ $LogLoss = -\frac{1}{N} \sum_{i=1}^N [y_i \log(p_i) - (1 - y_i) \log(1 - p_i)]$
with binary classification

Task Description

- ▶ We are given (training and testing) dataset /images, documents,.../
$$D = \{(x^{(l)}, y^{(l)})\}_{l \in L} \subset \mathbb{R}^n \times T$$
- ▶ The aim is to find a function $\hat{y} : \mathbb{R}^n \rightarrow T$ (target space), which estimates the testing data /unknown to the trained model/.
- ▶ In recommender systems (Online advertising) we deal with sparse input vectors $x^{(l)} \in \mathbb{R}^n$.
- ▶ T can be $\{-1, 1\}$, $\{0, 1\}$ (classification), \mathbb{R} (regression) or some categorical space (ranks for an item).

Task Description



Pipeline

- Each step generates many questions:
 - Data generation: **data types, sample size, online/offline...**
 - Preprocessing: **normalization, missing values, feature selection/extraction...**
 - Machine learning: **hypothesis, choice of learning paradigm/algorithm...**
 - Hypothesis validation: **cross-validation, model deployment...**

Ad Classification

- ▶ Classification from implicit feedback - the user does not rate explicitly..
- ▶ But we can "guess" items with which features he likes or does not care about.

Country	Day	Ad Type	Clicked ?
USA	3/3/15	MOVIE	1
China	1/7/14	GAME	0
China	3/3/15	GAME	1

Standard (dummy) encoding

USA	China	3/3/15	1/7/14	MOVIE	GAME	Clicked ?
1	0	1	0	1	0	1
0	1	0	1	0	1	0
0	1	1	0	0	1	1

- ▶ Very large feature space
- ▶ Very sparse samples (feature values)

(De)motivation!

Often features are more important in pairs:

"Country == USA" & "Day == Thanksgiving"

If we create a new feature (conjunction) for every pair of features..?!?

(De)motivation!

Often features are more important in pairs:

"Country == USA" & "Day == Thanksgiving"

If we create a new feature (conjunction) for every pair of features..?!?

- ▶ Feature space: insanely large

If originally n features \longrightarrow additional $\binom{n}{2} = \frac{n(n-1)}{2!}$

- ▶ Samples: still sparse

Problems with feature types

- Big number of features -> Dimensionality reduction -> SVD, PCA
 - **Dimensionality reduction:** “compress” the data from a high-dimensional representation into a lower-dimensional one (useful for visualization or as an internal transformation for other ML algorithms)
- Sparse features -> Hashing

SVD with ALS

Singular Value Decomposition (Matrix Factorization) with Alternate Least Squares.

The most basic matrix factorization model for recommender systems models the rating \hat{r} a user u would give to an item i by:

$$\hat{r}_{ui} = x_u^T y_i,$$


where

- ▶ $x_u^T = (x_u^1, \dots, x_u^N)$ - factor vector associated to the user
- ▶ $y_i^T = (y_i^1, \dots, y_i^N)$ - factor vector associated to the item
- ▶ The dimension N of the **factors** x_u^T, y_i^T is the **rank** of the model (factorization).


Factorization Models

Collaborative Filtering:

- Neighborhood Methods
- Latent Factor Methods



SVD – Dimensionality Reduction



- Instead of using two coordinates (x, y) to describe point locations, let's use only one coordinate (z)
- Point's position is its location along vector \mathbf{v}_1

SVD - Dimensionality Reduction

More details

- **Q:** How exactly is dim. reduction done?
- **A:** Set smallest singular values to zero

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 3 & 3 & 3 & 0 & 0 \\ 4 & 4 & 4 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 2 & 0 & 4 & 4 \\ 0 & 0 & 0 & 5 & 5 \\ 0 & 1 & 0 & 2 & 2 \end{bmatrix} \approx \begin{bmatrix} 0.13 & 0.02 & -0.01 \\ 0.41 & 0.07 & -0.03 \\ 0.55 & 0.09 & -0.04 \\ 0.68 & 0.11 & -0.05 \\ 0.15 & -0.59 & 0.65 \\ 0.07 & -0.73 & -0.67 \\ 0.07 & -0.29 & 0.32 \end{bmatrix} \times \begin{bmatrix} 12.4 & 0 & 0 \\ 0 & 9.5 & 0 \\ 0 & 0 & \cancel{1.3} \end{bmatrix} \times \begin{bmatrix} 0.56 & 0.59 & 0.56 & 0.09 & 0.09 \\ 0.12 & -0.02 & 0.12 & -0.69 & -0.69 \\ 0.40 & -0.80 & 0.40 & 0.09 & -0.09 \end{bmatrix}$$

SVD - Dimensionality Reduction

More details

- **Q:** How exactly is dim. reduction done?
- **A:** Set smallest singular values to zero

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 3 & 3 & 3 & 0 & 0 \\ 4 & 4 & 4 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 2 & 0 & 4 & 4 \\ 0 & 0 & 0 & 5 & 5 \\ 0 & 1 & 0 & 2 & 2 \end{bmatrix} \approx \begin{bmatrix} 0.13 & 0.02 & -0.01 \\ 0.41 & 0.07 & -0.03 \\ 0.55 & 0.09 & -0.04 \\ 0.68 & 0.11 & -0.05 \\ 0.15 & -0.59 & 0.65 \\ 0.07 & -0.73 & -0.67 \\ 0.07 & -0.29 & 0.32 \end{bmatrix} \times \begin{bmatrix} 12.4 & 0 & 0 \\ 0 & 9.5 & 0 \\ 0 & 0 & 1.3 \end{bmatrix} \times \begin{bmatrix} 0.56 & 0.59 & 0.56 & 0.09 & 0.09 \\ 0.12 & -0.02 & 0.12 & -0.69 & -0.69 \\ 0.40 & -0.80 & 0.40 & 0.09 & 0.09 \end{bmatrix}$$

The diagram illustrates the SVD decomposition of a matrix. The first matrix is a 7x5 matrix. The second matrix is a 7x3 matrix of singular values, with the third column (containing 1.3) crossed out with a red 'X'. The third matrix is a 3x5 matrix of right singular vectors, with the third row (containing 0.40, -0.80, 0.40, 0.09, 0.09) crossed out with a red line. The overall equation shows the first matrix approximated by the product of the second and third matrices.

Advantages of Factorization Models

Factorization models:

- ▶ Factorization of higher order interactions
- ▶ Efficient parameter estimation and superior performance
- ▶ Even for sparse data where just a few or no observations for those higher order effects are available

Main applications:

- ▶ Collaborative Filtering (in Recommender Systems)
- ▶ Link Prediction (in Social Networks)

Disadvantages of Factorization Models

BUT are not general prediction models. There are many factorization models designed for specific tasks.

- ▶ Standard models:
 - ▶ Matrix Factorization (MF)
 - ▶ Tensor Factorization (TF)
- ▶ Specific tasks:
 - ▶ SVD++ [Koren, Bell - Advances in CF]
 - ▶ timeSVD++ [Koren, Bell - Advances in CF]
 - ▶ PIFT (Pairwise Interaction Tensor Factorization)
 - ▶ FPMC (Factorizing Personalized Markov Chains)
 - ▶ And others...

Advantages of FMs

- ▶ FMs allow parameter estimation under very sparse data where SVMs fail.
- ▶ FMs have linear (and even "almost constant"!) complexity and don't rely on support vectors like SVMs.
- ▶ FMs are a general predictor that can work with any real valued feature vector. In contrast, other state-of-the-art factorization models work only on very restricted input data.
- ▶ Through suitable feature engineering, FMs can mimic state-of-the-art models like MF, SVD++, timeSVD++, PARAFAC, PITF, FPMC and others.

Model

Feature vector \mathbf{x}																	Target y					
$\mathbf{x}^{(1)}$	1	0	0	...	1	0	0	0	...	0.3	0.3	0.3	0	...	13	0	0	0	0	...	5	$y^{(1)}$
$\mathbf{x}^{(2)}$	1	0	0	...	0	1	0	0	...	0.3	0.3	0.3	0	...	14	1	0	0	0	...	3	$y^{(2)}$
$\mathbf{x}^{(3)}$	1	0	0	...	0	0	1	0	...	0.3	0.3	0.3	0	...	16	0	1	0	0	...	1	$y^{(2)}$
$\mathbf{x}^{(4)}$	0	1	0	...	0	0	1	0	...	0	0	0.5	0.5	...	5	0	0	0	0	...	4	$y^{(3)}$
$\mathbf{x}^{(5)}$	0	1	0	...	0	0	0	1	...	0	0	0.5	0.5	...	8	0	0	1	0	...	5	$y^{(4)}$
$\mathbf{x}^{(6)}$	0	0	1	...	1	0	0	0	...	0.5	0	0.5	0	...	9	0	0	0	0	...	1	$y^{(5)}$
$\mathbf{x}^{(7)}$	0	0	1	...	0	0	1	0	...	0.5	0	0.5	0	...	12	1	0	0	0	...	5	$y^{(6)}$
	A	B	C	...	TI	NH	SW	ST	...	TI	NH	SW	ST	...	Time	TI	NH	SW	ST	...		
	User				Movie					Other Movies rated					Time	Last Movie rated						

Notations

- ▶ For any $x \in \mathbb{R}^n$ let us define:
 - ▶ $m(x)$ to be the number of the nonzero elements of the feature vector;
 - ▶ \bar{m}_D to be the average of $m(x^{(l)})$ for all $l \in L$;
- ▶ Since we deal with huge sparsity, we have that $\bar{m}_D \ll n$ - the dimensionality of the feature space.
- ▶ One reason for huge sparsity is that the underlying problem deals with large categorical variable domains.

Definition of Factorization Machine Model of degree 2

$$\hat{y}(x) := w_0 + \sum_{i=1}^n w_i x_i + \sum_{i=1}^n \sum_{j=i+1}^n \langle v_i, v_j \rangle x_i x_j \quad (1)$$

where the model parameters that have to be estimated are:

- ▶ $w_0 \in \mathbb{R}$ – global bias (mean rating over all items/users)
- ▶ $w = (w_1, \dots, w_n) \in \mathbb{R}^n$ – weights of the corresponding features
- ▶ $v_i := (v_{i,1}, \dots, v_{i,k})$ within $V \in \mathbb{R}^{n \times k}$ describes the i -th variable(feature) with k factors.

Complexity of the model equation

- The model equation $\hat{y}(x)$ for every $x \in D$ can be computed in linear time $\mathcal{O}(kn)$ and under sparsity $\mathcal{O}(km(x))$

Proof:

$$\begin{aligned} \sum_{i=1}^n \sum_{j=i+1}^n \langle v_i, v_j \rangle x_i x_j &= \\ \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \langle v_i, v_j \rangle x_i x_j - \frac{1}{2} \sum_{i=1}^n \langle v_i, v_i \rangle x_i x_i &= \\ \frac{1}{2} \left[\sum_{i=1}^n \sum_{j=1}^n \sum_{f=1}^k v_{i,f} v_{j,f} x_i x_j - \sum_{i=1}^n \sum_{f=1}^k v_{i,f} v_{i,f} x_i x_i \right] &= \\ \frac{1}{2} \sum_{f=1}^k \left[\left(\sum_{i=1}^n v_{i,f} x_i \right) \left(\sum_{j=1}^n v_{j,f} x_j \right) - \sum_{i=1}^n v_{i,f}^2 x_i^2 \right] &= \\ \frac{1}{2} \sum_{f=1}^k \left[\left(\sum_{i=1}^n v_{i,f} x_i \right)^2 - \sum_{i=1}^n v_{i,f}^2 x_i^2 \right] \end{aligned}$$

Complexity of parameter updates

The needed computation to get the partial derivatives of $\hat{y}(x)$ is done while computing the computation of $\hat{y}(x)$. Therefore,

- ▶ All parameter updates for a case (x, y) can be done in $\mathcal{O}(kn)$ and under sparsity $\mathcal{O}(km(x))$.

$$\frac{\partial \hat{y}(x)}{\partial \theta} = \begin{cases} 1, & \text{if } \theta \text{ is } w_0 \\ x_i, & \text{if } \theta \text{ is } w_i \\ x_i \sum_{j=1}^n v_{j,f} x_j - v_{i,f} x_i^2 & \text{if } \theta \text{ is } v_{i,f} \end{cases} \quad (2)$$

General applications of FM

FMs can be applied to a variety of prediction tasks.

- ▶ **Regression:** $\hat{y}(x)$ can be directly applied as a regression predictor
- ▶ **Binary classification:** The sign of $\hat{y}(x)$ is used and the parameters are optimized for hinge loss or logit loss.
- ▶ **Ranking:** Input vectors x are ordered by the score of $\hat{y}(x)$ and optimization is done over pairs of instance vectors with a pairwise classification loss ([Joachims - Ranking CTR]).

Inference

Inference algorithms:

- ▶ SGD (Stochastic Gradient Descent)
- ▶ ALS (Alternate Least Squares)
- ▶ MCMC (Monte Carlo Markov Chains)

Library: libFM (C++)

- ▶ Author: Steffen Rendle
- ▶ Web site: <http://www.libfm.org/>
(very poor description of the library)
- ▶ Source code (C++): <https://github.com/srendle/libfm>
Very professionally written source code! (Linux and MacOS X)
- ▶ libFM features
 - ▶ stochastic gradient descent (SGD)
 - ▶ alternating least squares (ALS)
 - ▶ Bayesian inference using Markov Chain Monte Carlo (MCMC)

Library: fastFM (Python)

- ▶ Author: Immanuel Bayer
- ▶ Web site: <http://ibayer.github.io/fastFM/>
Very detailed tutorial for the usage of the library!
- ▶ Source code (Python):
<https://github.com/ibayer/fastFM>
Python (2.7 and 3.5) with the well known scikit-learn API.
Performance critical code in C and wrapped with Cython.

Task	Solver	Loss
Regression	ALS, SGD, MCMC	Square Loss
Classification	ALS, SGD, MCMC	Probit(MAP), Probit, Sigmoid
Ranking	SGD	[Rendle - BPR]

Library: LIBFFM (C++)






A library for Field-aware Factorization Machines from the Machine Learning group at National Taiwan University

- ▶ Author: YuChin Juan and colleagues
- ▶ Web site:
<https://www.csie.ntu.edu.tw/~cjlin/libffm/>
(very poor description of the library)
- ▶ Source code (C++):
<https://github.com/guestwalk/libffm>
- ▶ FFM is prone to overfitting and not very stable for the moment.

Table of contents

Thank you for your attention!

References

-  Steffen Rendle *Factorization Machines*. 2010. <http://www.ismll.uni-hildesheim.de/pub/pdfs/Rendle2010FM.pdf>
-  Steffen Rendle, *Christoph Freudenthaler, Zeno Gantner, and Lars Schmidt-Thieme*. *Bpr: Bayesian personalized ranking from implicit feedback*. *UAI '09*, pages 452 2009
-  Koren Y. and Bell R. *Advances in Collaborative Filtering*.
-  Blondel M. *Convex Factorization Machines*.
<http://www.mblondel.org/publications/mblondel-ecmlpkdd2015.pdf>
-  Thorsten Joachims, *Optimizing Search Engines using Clickthrough Data* 2002.